

Traffic_Sign_Classifier

October 2, 2018

1 Self-Driving Car Engineer Nanodegree

1.1 Deep Learning

1.2 Project: Build a Traffic Sign Recognition Classifier

1.3 Author: Felipe Pamboukian

In this notebook, a traffic sign classifier is implemented using Convolutional Neural Network based on LeNet architecture, German Traffic Sign Dataset is used to train the model.

1.4 Step 0: Load The Data

```
In [1]: import cv2
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
from sklearn.utils import shuffle
import tensorflow as tf
from tensorflow.contrib.layers import flatten
```

```
In [2]: # Load pickled data
import pickle

training_file = '../data/train.p'
validation_file= '../data/valid.p'
testing_file = '../data/test.p'

with open(training_file, mode='rb') as f:
    train = pickle.load(f)
with open(validation_file, mode='rb') as f:
    valid = pickle.load(f)
with open(testing_file, mode='rb') as f:
    test = pickle.load(f)
```

```
X_train, y_train = train['features'], train['labels']
X_valid, y_valid = valid['features'], valid['labels']
X_test, y_test = test['features'], test['labels']
```

1.5 Step 1: Dataset Summary & Exploration

The pickled data is a dictionary with 4 key/value pairs:

- 'features' is a 4D array containing raw pixel data of the traffic sign images, (num examples, width, height, channels).
- 'labels' is a 1D array containing the label/class id of the traffic sign. The file `signnames.csv` contains id -> name mappings for each id.
- 'sizes' is a list containing tuples, (width, height) representing the original width and height the image.
- 'coords' is a list containing tuples, (x1, y1, x2, y2) representing coordinates of a bounding box around the sign in the image. **THESE COORDINATES ASSUME THE ORIGINAL IMAGE. THE PICKLED DATA CONTAINS RESIZED VERSIONS (32 by 32) OF THESE IMAGES**

1.5.1 Basic Summary of the Dataset

```
In [3]: n_train = len(X_train)
        n_validation = len(X_valid)
        n_test = len(X_test)

        image_shape = X_train[0].shape

        classes = pd.DataFrame(y_train)[0].unique()
        n_classes = len(classes)

        print("Number of training examples =", n_train)
        print("Number of testing examples =", n_test)
        print("Image data shape =", image_shape)
        print("Number of classes =", n_classes)
```

```
Number of training examples = 34799
Number of testing examples = 12630
Image data shape = (32, 32, 3)
Number of classes = 43
```

1.5.2 Visualization of the dataset

Showing a sample for each class

```
In [4]: ### Data exploration visualization code goes here.
        ### Feel free to use as many code cells as needed.
```

```

import matplotlib.cm as cm
# Visualizations will be shown in the notebook.
%matplotlib inline

def plot_each_class(x, y, classes):
    fig = plt.figure(figsize=(12,19))
    for j in range(len(classes)-1):
        ax = fig.add_subplot(15, 11, j+1)
        index = np.random.choice(np.where(y==j)[0])
        if x.shape[3] == 1:
            ax.imshow(x[index], cmap='gray')
        else:
            ax.imshow(x[index])
        plt.xticks(np.array([]))
        plt.yticks(np.array([]))
        plt.title('Class:' + str(j))
        plt.tight_layout()

plot_each_class(X_train, y_train, classes)

```



Showing number of samples for each class

```

In [5]: class_arr= []
        samples_arr=[]


        for class_n in range(n_classes):
            class_indices = np.where(y_train == class_n)
            n_samples = len(class_indices[0])
            class_arr.append(class_n)
            samples_arr.append(n_samples)

```

```

plt.hist(y_train, bins=43)
plt.bar(class_arr, samples_arr, align='center', alpha=0.5)
plt.ylabel('Classes')
plt.xlabel('No of Samples')
plt.title('Data Visualization')
plt.show()

```



1.6 Step 2: Design and Test a Model Architecture

Design and implement a deep learning model that learns to recognize traffic signs. Train and test your model on the [German Traffic Sign Dataset](#).

1.6.1 Pre-process the Data Set (normalization, grayscale, etc.)

I got better results using color images, so now I just normalize.

```

In [6]: def preprocessing(imageset):
        new_shape = image_shape[0:2] + (3,)
        new_imageset = np.empty(shape=(len(imageset),) + new_shape, dtype=int)

        for i in range(0, len(imageset)):
            #gray_img = cv2.cvtColor(imageset[i], cv2.COLOR_RGB2GRAY)

```

```

        #norm_img = cv2.normalize(gray_img, np.zeros(image_shape[0:2]), 0, 255, cv2.NORM
norm_img = cv2.normalize(imageset[i], np.zeros(image_shape[0:2]), 0, 255, cv2.NO
new_imageset[i] = np.reshape(norm_img, new_shape)

    return new_imageset

#X_train = preprocessing(X_train)
#X_valid = preprocessing(X_valid)
#X_test = preprocessing(X_test)

```

1.6.2 Shuffle training dataset

```
In [7]: X_train, y_train = shuffle(X_train, y_train)
```

1.6.3 Model Architecture

The model architecture is almost the same, I just modified to use color images and output to 43 classes.

```
In [8]: def LeNet(x, keep_prob):
    # Arguments used for tf.truncated_normal, randomly defines variables for the weights and
    mu = 0
    sigma = 0.1

    # Layer 1: Convolutional. Input = 32x32x3. Output = 28x28x6.
    conv1_W = tf.Variable(tf.truncated_normal(shape=(5, 5, 3, 6), mean = mu, stddev = sigma))
    conv1_b = tf.Variable(tf.zeros(6))
    conv1 = tf.nn.conv2d(x, conv1_W, strides=[1, 1, 1, 1], padding='VALID') + conv1_b
    # Activation.
    conv1 = tf.nn.relu(conv1)
    # Pooling. Input = 28x28x6. Output = 14x14x6.
    conv1 = tf.nn.max_pool(conv1, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='VALID')

    # Layer 2: Convolutional. Output = 10x10x16.
    conv2_W = tf.Variable(tf.truncated_normal(shape=(5, 5, 6, 16), mean = mu, stddev = sigma))
    conv2_b = tf.Variable(tf.zeros(16))
    conv2 = tf.nn.conv2d(conv1, conv2_W, strides=[1, 1, 1, 1], padding='VALID') + conv2_b
    # Activation.
    conv2 = tf.nn.relu(conv2)
    # Pooling. Input = 10x10x16. Output = 5x5x16.
    conv2 = tf.nn.max_pool(conv2, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='VALID')
    # Flatten. Input = 5x5x16. Output = 400.
    fc0 = flatten(conv2)
    fc0 = tf.nn.dropout(fc0, keep_prob=keep_prob)

    # Layer 3: Fully Connected. Input = 400. Output = 120.
    fc1_W = tf.Variable(tf.truncated_normal(shape=(400, 120), mean = mu, stddev = sigma))
    fc1_b = tf.Variable(tf.zeros(120))

```

```

fc1    = tf.matmul(fc0, fc1_W) + fc1_b
# Activation.
fc1     = tf.nn.relu(fc1)

# Layer 4: Fully Connected. Input = 120. Output = 84.
fc2_W  = tf.Variable(tf.truncated_normal(shape=(120, 84), mean = mu, stddev = sigma))
fc2_b  = tf.Variable(tf.zeros(84))
fc2    = tf.matmul(fc1, fc2_W) + fc2_b
# Activation.
fc2     = tf.nn.relu(fc2)

# Layer 5: Fully Connected. Input = 84. Output = 43.
fc3_W  = tf.Variable(tf.truncated_normal(shape=(84, 43), mean = mu, stddev = sigma))
fc3_b  = tf.Variable(tf.zeros(43))
logits = tf.matmul(fc2, fc3_W) + fc3_b

return logits

```

1.6.4 Train, Validate and Test the Model

A validation set can be used to assess how well the model is performing. A low accuracy on the training and validation sets imply underfitting. A high accuracy on the training set but low accuracy on the validation set implies overfitting.

```

In [9]: ### Train your model here.
### Calculate and report the accuracy on the training and validation set.
### Once a final model architecture is selected,
### the accuracy on the test set should be calculated and reported as well.
### Feel free to use as many code cells as needed.

x = tf.placeholder(tf.float32, (None, 32, 32, 3))
y = tf.placeholder(tf.int32, (None))
one_hot_y = tf.one_hot(y, 43)

EPOCHS = 50
BATCH_SIZE = 128
rate = 0.0009
KEEP_PROB = 0.5

keep_prob = tf.placeholder(tf.float32)

logits = LeNet(x, keep_prob)
cross_entropy = tf.nn.softmax_cross_entropy_with_logits(labels=one_hot_y, logits=logits)
loss_operation = tf.reduce_mean(cross_entropy)
optimizer = tf.train.AdamOptimizer(learning_rate = rate)
training_operation = optimizer.minimize(loss_operation)

correct_prediction = tf.equal(tf.argmax(logits, 1), tf.argmax(one_hot_y, 1))

```

```

accuracy_operation = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
saver = tf.train.Saver()

def evaluate(X_data, y_data):
    num_examples = len(X_data)
    total_accuracy = 0
    sess = tf.get_default_session()
    for offset in range(0, num_examples, BATCH_SIZE):
        batch_x, batch_y = X_data[offset:offset+BATCH_SIZE], y_data[offset:offset+BATCH_SIZE]
        accuracy = sess.run(accuracy_operation, feed_dict={x: batch_x, y: batch_y, keep_prob: 1.0})
        total_accuracy += (accuracy * len(batch_x))
    return total_accuracy / num_examples

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    num_examples = len(X_train)

    print("Training...")
    print()
    for i in range(EPOCHS):
        X_train, y_train = shuffle(X_train, y_train)
        for offset in range(0, num_examples, BATCH_SIZE):
            end = offset + BATCH_SIZE
            batch_x, batch_y = X_train[offset:end], y_train[offset:end]
            sess.run(training_operation, feed_dict={x: batch_x, y: batch_y, keep_prob: 0.5})

        validation_accuracy = evaluate(X_valid, y_valid)
        print("EPOCH {} ...".format(i+1))
        print("Validation Accuracy = {:.3f}".format(validation_accuracy))
        print()

    saver.save(sess, './lenet')
    print("Model saved")

```

Training...

EPOCH 1 ...

Validation Accuracy = 0.375

EPOCH 2 ...

Validation Accuracy = 0.598

EPOCH 3 ...

Validation Accuracy = 0.722

EPOCH 4 ...

Validation Accuracy = 0.806

EPOCH 5 ...
Validation Accuracy = 0.837

EPOCH 6 ...
Validation Accuracy = 0.861

EPOCH 7 ...
Validation Accuracy = 0.870

EPOCH 8 ...
Validation Accuracy = 0.875

EPOCH 9 ...
Validation Accuracy = 0.897

EPOCH 10 ...
Validation Accuracy = 0.904

EPOCH 11 ...
Validation Accuracy = 0.890

EPOCH 12 ...
Validation Accuracy = 0.900

EPOCH 13 ...
Validation Accuracy = 0.912

EPOCH 14 ...
Validation Accuracy = 0.911

EPOCH 15 ...
Validation Accuracy = 0.921

EPOCH 16 ...
Validation Accuracy = 0.920

EPOCH 17 ...
Validation Accuracy = 0.930

EPOCH 18 ...
Validation Accuracy = 0.913

EPOCH 19 ...
Validation Accuracy = 0.916

EPOCH 20 ...
Validation Accuracy = 0.926

EPOCH 21 ...
Validation Accuracy = 0.931

EPOCH 22 ...
Validation Accuracy = 0.927

EPOCH 23 ...
Validation Accuracy = 0.931

EPOCH 24 ...
Validation Accuracy = 0.936

EPOCH 25 ...
Validation Accuracy = 0.936

EPOCH 26 ...
Validation Accuracy = 0.933

EPOCH 27 ...
Validation Accuracy = 0.926

EPOCH 28 ...
Validation Accuracy = 0.937

EPOCH 29 ...
Validation Accuracy = 0.935

EPOCH 30 ...
Validation Accuracy = 0.942

EPOCH 31 ...
Validation Accuracy = 0.933

EPOCH 32 ...
Validation Accuracy = 0.938

EPOCH 33 ...
Validation Accuracy = 0.929

EPOCH 34 ...
Validation Accuracy = 0.936

EPOCH 35 ...
Validation Accuracy = 0.937

EPOCH 36 ...
Validation Accuracy = 0.941

EPOCH 37 ...
Validation Accuracy = 0.942

EPOCH 38 ...
Validation Accuracy = 0.932

EPOCH 39 ...
Validation Accuracy = 0.949

EPOCH 40 ...
Validation Accuracy = 0.941

EPOCH 41 ...
Validation Accuracy = 0.944

EPOCH 42 ...
Validation Accuracy = 0.941

EPOCH 43 ...
Validation Accuracy = 0.944

EPOCH 44 ...
Validation Accuracy = 0.943

EPOCH 45 ...
Validation Accuracy = 0.944

EPOCH 46 ...
Validation Accuracy = 0.941

EPOCH 47 ...
Validation Accuracy = 0.944

EPOCH 48 ...
Validation Accuracy = 0.946

EPOCH 49 ...
Validation Accuracy = 0.939

EPOCH 50 ...
Validation Accuracy = 0.942

Model saved

1.6.5 Evaluate trained model using testset

```
In [16]: with tf.Session() as sess:
          saver.restore(sess, tf.train.latest_checkpoint('.'))

          test_accuracy = evaluate(X_test, y_test)
          print("Test Accuracy = {:.3f}".format(test_accuracy))

INFO:tensorflow:Restoring parameters from ./lenet
Test Accuracy = 0.942
```

1.7 Step 3: Test a Model on New Images

Now I applied the trained model to the German traffic sign images that were obtained from the Internet.

1.7.1 Load and Output the Images

```
In [22]: ### Load the images and plot them here.
          ### Feel free to use as many code cells as needed.
          import glob

          test_signs = glob.glob('./test_images/*')
          X_signs = []

          sign_names = pd.read_csv('signnames.csv')
          y_sign_names = sign_names['SignName']

          #y_signs = [25,40,14,3,11] # CPU
          y_signs = [14,40,3,11,25] # GPU

          for i in range(len(test_signs)):
              img = mpimg.imread(test_signs[i])
              img = cv2.resize(img, (32,32))
              X_signs.append(img)

          def plot_test_signs(X_signs):
              plt.figure(figsize=(10,5))
              plt.subplot(231)
              plt.title('Image #1, Class: ' + str(y_signs[0]))
              plt.imshow(X_signs[0])
              plt.subplot(232)
              plt.title('Image #2, Class: ' + str(y_signs[1]))
              plt.imshow(X_signs[1])
              plt.subplot(233)
              plt.title('Image #3, Class: ' + str(y_signs[2]))
```

```

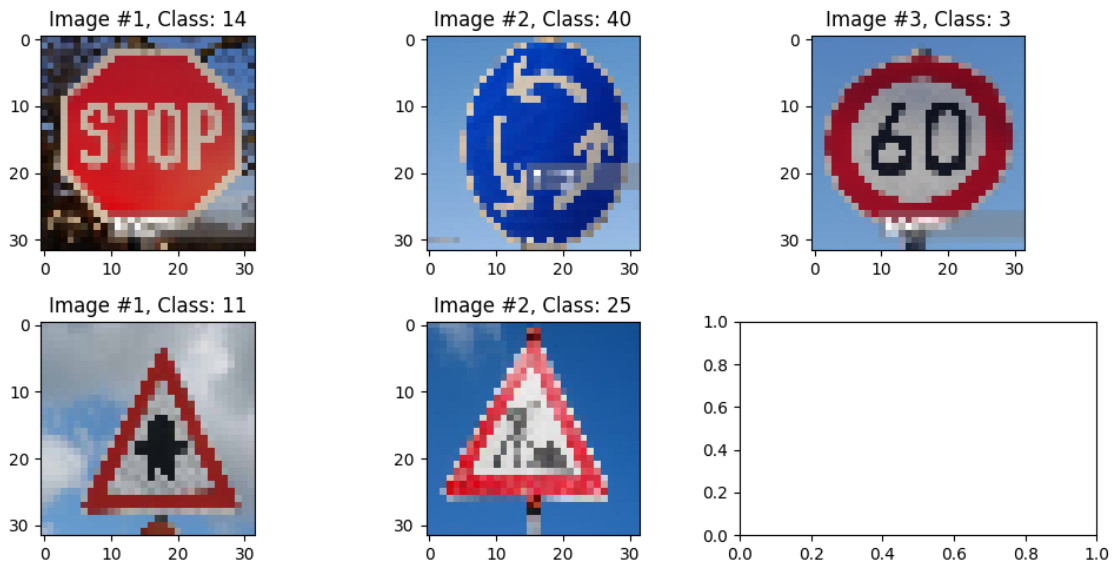
plt.imshow(X_signs[2])
plt.subplot(234)
plt.title('Image #1, Class: ' + str(y_signs[3]))
plt.imshow(X_signs[3])
plt.subplot(235)
plt.title('Image #2, Class: ' + str(y_signs[4]))
plt.imshow(X_signs[4])
plt.subplot(236)
plt.tight_layout()
plt.show()
return None

```

```

plot_test_signs(X_signs)
X_signs_new = X_signs
#X_signs_new = preprocessing(X_signs)

```



1.7.2 Predict the Sign Type for Each Image

```

In [23]: ### Run the predictions here and use the model to output the prediction for each image.
### Make sure to pre-process the images with the same pre-processing pipeline used early.
### Feel free to use as many code cells as needed.
with tf.Session() as sess:
    saver.restore(sess, tf.train.latest_checkpoint('.'))
    prediction = tf.argmax(logits, 1)
    test_prediction = sess.run(
        prediction,
        feed_dict={x: X_signs_new, y: y_signs, keep_prob: 1})

```

```

print (test_prediction)
print ()
for i in range(len(test_prediction)):
    print ('Prediction:', y_sign_names[test_prediction[i]], test_prediction[i])

```

INFO:tensorflow:Restoring parameters from ./lenet

[14 40 3 11 22]

Prediction: Stop 14
Prediction: Roundabout mandatory 40
Prediction: Speed limit (60km/h) 3
Prediction: Right-of-way at the next intersection 11
Prediction: Bumpy road 22

1.7.3 Analyze Performance

```

In [24]: ### Calculate the accuracy for these 5 new images.
        ### For example, if the model predicted 1 out of 5 signs correctly, it's 20% accurate
        with tf.Session() as sess:
            saver.restore(sess, tf.train.latest_checkpoint('.'))

            test_accuracy = evaluate(X_signs_new, y_signs)
            print("Test Accuracy = {:.3f}".format(test_accuracy))

```

INFO:tensorflow:Restoring parameters from ./lenet

Test Accuracy = 0.800

1.7.4 Output Top 5 Softmax Probabilities For Each Image Found on the Web

```

In [26]: # Print out the top five softmax probabilities for the predictions on
        # the German traffic sign images found on the web.
        softmax_logits = tf.nn.softmax(logits)
        top_k = tf.nn.top_k(softmax_logits, k=5)

        with tf.Session() as sess:
            sess.run(tf.global_variables_initializer())
            saver.restore(sess, tf.train.latest_checkpoint('.'))
            my_softmax_logits = sess.run(softmax_logits, feed_dict={x: X_signs_new, keep_prob: 1.0})
            my_top_k = sess.run(top_k, feed_dict={x: X_signs_new, keep_prob: 1.0})
            print(my_softmax_logits)
            print()
            print(my_top_k)

```

INFO:tensorflow:Restoring parameters from ./lenet

8.33419095e-14	2.47560184e-11	4.46718671e-14	8.96737025e-12
6.62408428e-14	2.11653442e-14	1.42625840e-21	5.31765431e-16
3.30470624e-15	6.24675225e-13	8.58986710e-13	2.09475297e-12

3.11439188e-13	4.95133830e-11	9.99997020e-01	6.76786406e-17
8.96636384e-19	2.93159883e-06	3.64587804e-11	3.89979329e-17
6.42612345e-15	1.51237002e-15	6.13543895e-12	2.17136557e-13
4.27939849e-15	7.07075642e-08	2.31084107e-14	7.61308033e-19
2.25605897e-13	1.64613687e-11	2.70337797e-11	3.22322963e-12
4.35938206e-20	2.45900803e-23	1.82881655e-15	7.30913253e-16
4.95172539e-15	5.70789115e-21	7.33046559e-16	1.86075662e-24
2.25361791e-19	1.14734980e-17	1.00270483e-16]	
[8.64663582e-28	6.14643158e-16	2.54786169e-20	9.98630656e-25
2.08731215e-24	7.23404746e-21	1.29246390e-29	1.24316172e-17
6.15030596e-22	9.98079944e-22	3.52489997e-25	5.60277793e-23
1.85616855e-16	1.09113196e-22	2.43387561e-29	3.83000675e-27
1.91021207e-20	3.73936191e-25	2.01966871e-24	4.15749944e-33
3.07489953e-22	1.05950905e-27	7.87922751e-34	5.13492489e-25
3.53437727e-26	5.66590891e-21	5.89116376e-21	3.52302487e-36
4.93696804e-24	4.13014945e-27	1.18216484e-29	2.50350752e-25
4.18412455e-22	6.43445783e-15	4.51513136e-16	1.37925490e-10
2.28020328e-11	3.39953592e-18	3.81641902e-13	3.40208288e-23
1.00000000e+00	7.94834118e-16	2.97009442e-20]	
[6.82987795e-22	2.48193386e-16	1.21937174e-11	1.00000000e+00
4.45892236e-23	1.69297557e-18	4.94503767e-38	6.97144847e-28
5.61666881e-37	8.30827548e-23	6.21899111e-27	3.57340925e-25
3.29425906e-36	3.44215257e-27	2.75572611e-27	3.60383574e-28
2.06479357e-35	2.48102584e-38	3.86937381e-29	4.77271442e-31
3.43523498e-37	2.14230861e-20	1.69024039e-32	1.52697786e-22
1.36462585e-27	1.15915675e-24	3.48185146e-34	7.24856425e-32
2.94186528e-31	1.78283142e-27	5.04649795e-22	8.97375901e-24
1.06699886e-30	1.26347658e-31	3.31033413e-28	5.87372451e-22
2.36415676e-35	7.10756747e-35	3.05792248e-32	0.00000000e+00
3.76908038e-24	3.33743136e-35	2.81900994e-35]	
[9.41088257e-27	6.76190860e-25	3.37024918e-23	9.97344024e-21
1.88880859e-30	9.90461094e-24	3.13453427e-38	4.16800970e-23
7.44358758e-30	5.50617834e-28	4.14711169e-28	9.96304750e-01
5.26572322e-23	1.16020399e-30	4.85261287e-34	7.50643383e-30
1.49179799e-27	1.51698119e-32	3.43354477e-11	1.55815430e-16
7.29498807e-24	7.03804848e-10	4.71729281e-34	4.46725212e-10
1.96114291e-10	2.51008394e-16	1.46678565e-16	1.13317238e-11
6.48777690e-13	4.00260225e-17	3.69528192e-03	1.52550584e-17
5.13945473e-35	8.49671267e-33	1.16259878e-26	1.29697429e-31
0.00000000e+00	0.00000000e+00	1.34675659e-26	0.00000000e+00
4.48167510e-25	1.58831045e-31	1.81321555e-31]	
[1.98316813e-13	3.19795218e-10	2.49086899e-16	1.47970438e-11
3.38905228e-12	3.29701752e-11	2.42485169e-20	1.49230061e-14
4.48127592e-13	1.69701251e-16	2.59623634e-09	2.97087945e-11
2.20081183e-15	7.81289801e-13	4.83160678e-10	6.59537147e-11
5.94102954e-17	1.19752944e-10	8.56533866e-09	1.75362786e-17
1.03315810e-14	1.05865788e-10	9.99954462e-01	3.61851933e-14
6.46400996e-14	4.43562894e-05	1.05863876e-06	9.51114746e-15